NLP for Finance

Cris Chai, Dhruva Iyer, Dristi Roy, Sean Wu, Legolas Zhang, Satvik Marthi



We used NLP to predict sentiments of stock market tweets, involving preprocessing and inputting them to train classification models.

Sentiments can help predict changes in stock prices.

NLP pre-processing Simplify the data **Tokenizatio** Split the sentence The stock will go up

will

go

up

stock

The



Text Cleaning

- Text cleaning is the process of simplifying a sentence
- One way is by replacing characters such as , ! ; " '
 with spaces
- Sentences will have the same meaning

Unsimplified Sentence: I was surprised when the stock price increased.

Simplified Sentence: The stock price increased.

Stopword removal

The company is not good

The company is good

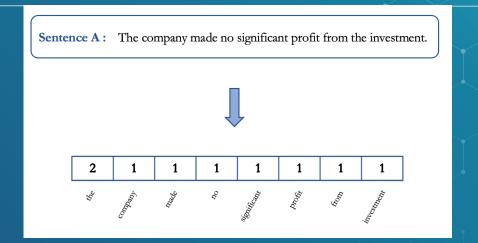
Remove stopwords like

- there
- ♦ is
- the
- and

company good

Vectorization

Converts a list of words into a vector of numbers



Data set balancing

- Data is not balanced
- Lots of neutral data

We did not address this, but we should if we do this again.



0 is negative1 is neutral2 is positive

One Hot Encoding

One hot encoding is converting categorical data into numerical data

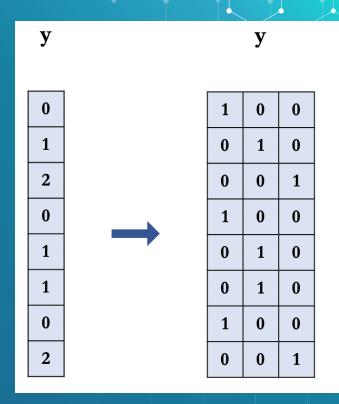
- 1 column per possible label
- Easy to interpret model's output

Before:

0: negative, 1: neutral, 2: positive

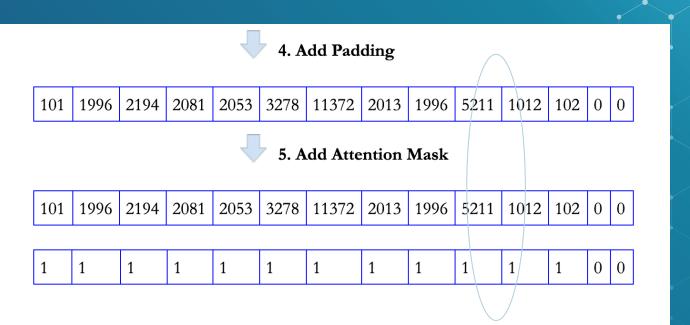
After:

Col 1: negative, Col 2: neutral, Col 3: positive



Padding and Attention Masks

- Extends vectors with padding to have equal length inputs
- Attention mask draws model attention away from padding



Word Order

- Changes the meaning of a sentence
- Very important

Sentence A: The company made no significant profit from the investment.

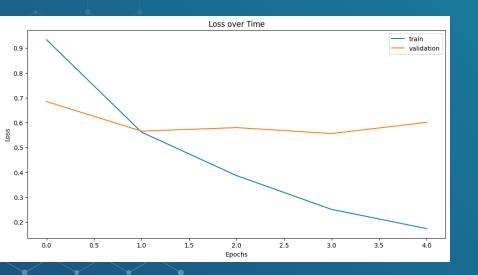
Sentence B: No, the company made significant profit from the investment.

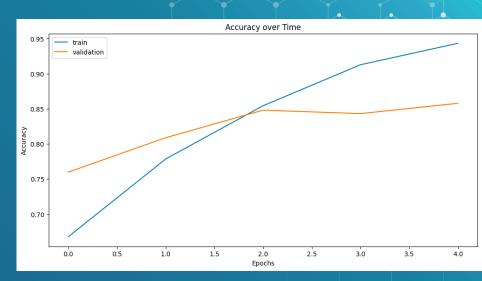
LSTM

- LSTM stands for long short-term memory
- It can help to solve word order issues
- A sequence of words rather than bag



LSTM Result



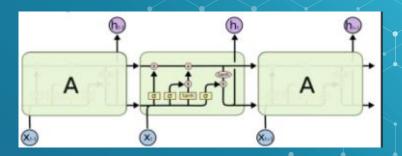


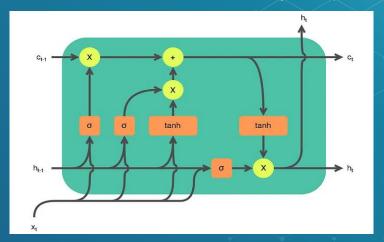
Model Architecture and Hyperparameters

Model layers Hyperparameters

- Embedding Epochs
- Dropout
- LSTM
- Dense

- Max Sequence Length
- Max NB Words







BERT is a pre-trained model from Google

- Very accurate
- Needs less effort than new model
- Speeds up training
- Needs less data



Setting Up BERT

Pre-processing steps

- Add special tokens
- Tokenize
- Convert tokens to indexes

The company made no significant profit from the investment.

1. Add Special Tokens

"[CLS] The company made no significant profit from the investment. [SEP]"

2. Tokenize

[CLS] the company made no significant profit from the investment . [SEP]

3. Map Tokens to Indices

101 1996 2194 2081 2053 3278 11372 2013 1996 5211 1012 102

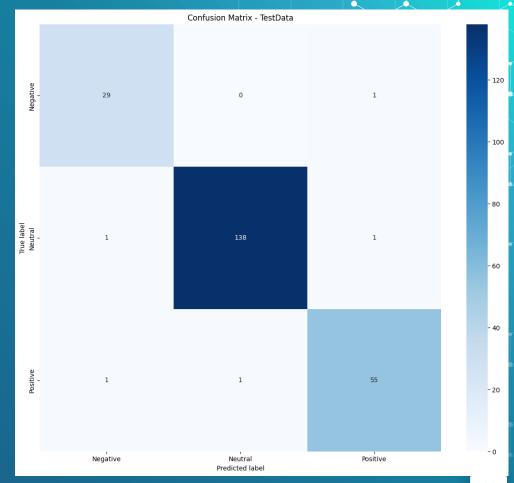
Add padding and attention mask

Using the model

- Huggingface Pytorch has a BERT implementation
- Add an output layer to BERT
- Train like any other model

BERT Results

Highly accurate in testing with no hyperparameter optimization Only 5/227 wrong answers (98% accuracy)



Thank You